

Prepare data for machine learning

ANALYZING IOT DATA IN PYTHON



Matthias Voppichler
IT Developer

Machine Learning Refresher

- Supervised learning
 - Classification
 - Regression
- Unsupervised learning
 - Cluster analysis
- Deep learning
 - Neural networks

Machine Learning Refresher

- Supervised learning
 - **Classification**
 - Regression
- Unsupervised learning
 - Cluster analysis
- Deep learning
 - Neural networks

Labels

```
print(environment_labeled.head())
```

		humidity	temperature	pressure	label
timestamp					
2018-10-01	00:00:00	81.0	11.8	1013.4	1
2018-10-01	00:15:00	79.7	11.9	1013.1	1
2018-10-01	00:30:00	81.0	12.1	1013.0	1
2018-10-01	00:45:00	79.7	11.7	1012.7	1
2018-10-01	01:00:00	84.3	11.2	1012.6	1

Train / Test split


Splitting time series data

- Model should not see test-data during training
- Cannot use random split
- Model should not be allowed to look into the future

Train / test split

```
split_day = "2018-10-13"  
train = environment[:split_day]  
test = environment[split_day:]  
print(train.iloc[0].name)  
print(train.iloc[-1].name)  
print(test.iloc[0].name)  
print(test.iloc[-1].name)
```

```
2018-10-01 00:00:00  
2018-10-13 23:45:00  
2018-10-13 00:00:00  
2018-10-15 23:45:00
```



Features and Labels

```
X_train = train.drop("target", axis=1)
y_train = train["target"]
X_test = test.drop("target", axis=1)
y_test = test["target"]
print(X_train.shape)
print(y_train.shape)
```

```
(1248, 3)
(1248,)
```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression  
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)  
print(logreg.predict(X_test))
```

```
[0 0 1 1 1 1 1 0 0]
```


Let's practice!

ANALYZING IOT DATA IN PYTHON

Scaling data for machine learning

ANALYZING IOT DATA IN PYTHON



Matthias Voppichler
IT Developer

Evaluate the model

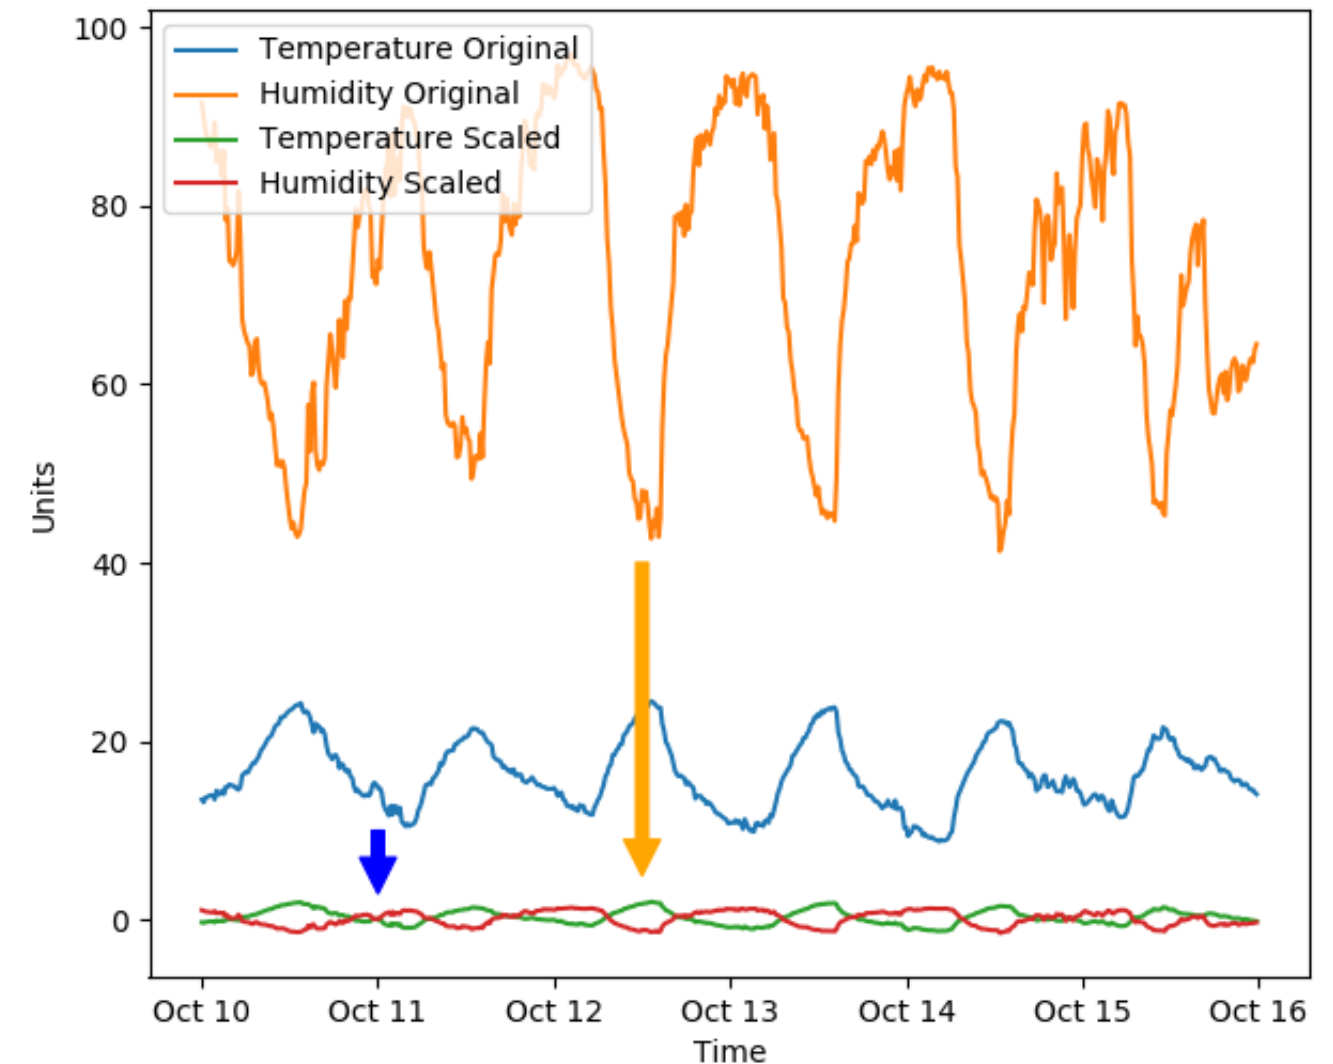
```
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)  
print(logreg.score(X_test, y_test))
```

```
0.78145113
```

Scaling

scikit-learn's `StandardScaler`

- remove mean
- scale data to variance



Unscaled data

```
print(data.head())
```

	humidity	temperature	pressure
timestamp			
2018-10-01 00:00:00	81.0	11.8	1013.4
2018-10-01 00:15:00	79.7	11.9	1013.1
2018-10-01 00:30:00	81.0	12.1	1013.0
2018-10-01 00:45:00	79.7	11.7	1012.7
2018-10-01 01:00:00	84.3	11.2	1012.6

StandardScaler

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
sc.fit(data)
print(sc.mean_)
print(sc.var_)
```

```
[ 71.8826716    14.17002019 1018.17042396]
[372.78261022  20.37926608  53.67519188]
```

```
data_scaled = sc.transform(data)
```

StandardScaler

```
df_scaled = pd.DataFrame(data_scaled,  
                           columns=data.columns,  
                           index=data.index)  
  
print(data_scaled.head())
```

	humidity	temperature	pressure
timestamp			
2018-10-01 00:00:00	0.472215	-0.524998	-0.651134
2018-10-01 00:15:00	0.404884	-0.502847	-0.692082
2018-10-01 00:30:00	0.472215	-0.458543	-0.705731
2018-10-01 00:45:00	0.404884	-0.547150	-0.746679
2018-10-01 01:00:00	0.643132	-0.657908	-0.760329

Evaluate the model

```
logreg = LogisticRegression()  
logreg.fit(X_train_scaled, y_train_scaled)  
  
print(logreg.score(X_test_scaled, y_test_scaled))
```

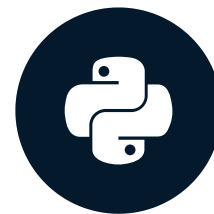
```
0.88145113
```


Let's practice!

ANALYZING IOT DATA IN PYTHON

Develop machine learning pipeline

ANALYZING IOT DATA IN PYTHON



Matthias Voppichler
IT Developer

Pipeline

- Transform
 - Conversation
 - Scaling
- Estimator
 - Model

Create a Pipeline

```
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
# Initialize Objects
sc = StandardScaler()
logreg = LogisticRegression()
# Create pipeline
pl = Pipeline([
    ("scale", sc),
    ("logreg", logreg)
])
```

Inspect Pipeline

```
pipeline
```

```
Pipeline(memory=None,  
          steps=[('scale', StandardScaler(copy=True, with_mean=True, with_std=True)),  
                 ('logreg', <class 'sklearn.linear_model.logistic.LogisticRegression'>)])
```

```
pipeline.fit(X_train, y_train)  
print(pipeline.predict(X_test))
```

```
[0 0 1 1 0 1 1 0 0]
```

Save model

```
import pickle
with Path("pipeline_model.pkl").open("bw") as f:
    pickle.dump(pl, f)
```

Load Model

```
import pickle
with Path("pipeline_model.pkl").open('br') as f:
    pl = pickle.load(f)
```

pl

```
Pipeline(memory=None,
          steps=[('scale', StandardScaler(copy=True, with_mean=True, with_std=True)),
                 ('logreg', LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                 intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l2',
                 random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False))])
```

A word of caution

DO NOT unpickle untrusted files, this can lead to malicious code being executed.

Let's practice!

ANALYZING IOT DATA IN PYTHON

Apply a machine learning model

ANALYZING IOT DATA IN PYTHON



Matthias Voppichler
IT Developer

Model Recap

```
# Create Pipeline
pl = Pipeline([
    ("scale", StandardScaler()),
    ("logreg", LogisticRegression())
])

# Fit the pipeline
pl.fit(X_train, y_train)
print(pl.score(X_test, y_test))
```

```
0.8897932222860425
```

Predict

```
predictions = pl.predict(X_test)
print(predictions)
print(f"Test length: {len(X_test)}")
print(f"Prediction length: {len(predictions)}")
```

```
[0 0 0 ... 1 1 1]
Test length: 500
Prediction length: 500
```

Record conversation

```
print(single_record)
```

```
{'timestamp': '2018-11-30 18:15:00',  
  'humidity': 81.7,  
  'pressure': 1019.8,  
  'temperature': 1.5},
```

```
cols = X_train.columns  
df = pd.DataFrame.from_records([single_record],  
                               index="timestamp",  
                               columns=cols)
```

Apply to datastream


```
def on_message(client, userdata, message):  
    data = json.loads(message.payload)  
    df = pd.DataFrame.from_records([data],  
                                   index="timestamp",  
                                   columns=cols)  
  
    category = pl.predict(df)  
    maybe_alert(category[0])  
subscribe.callback(on_message, topic, hostname=MQTT_HOST)
```

Let's practice!

ANALYZING IOT DATA IN PYTHON

Wrapping up

ANALYZING IOT DATA IN PYTHON



Matthias Voppichler
IT Developer

What you have learned

- Accessing IoT data
 - from a REST API
 - from a datastream
- Data Cleaning
- Correlations
- Time series decomposition
- Machine learning pipeline

Next steps

- Machine Learning
- Database
- Big data
- PySpark

Congratulations!

ANALYZING IOT DATA IN PYTHON